

MongoDB Mastery



Content

- Relationship & Advanced Data Modeling
- Aggregation Pipeline
- Indexing for Performance
- Database Transactions

Relationship Modeling



Relationship Modeling



- **What Is Relationship Modeling?**
 - In relational databases, relationships are defined via foreign keys
 - In MongoDB, relationships are created using document structure
 - Embedded documents or referenced documents
 - MongoDB is flexible — choose what's best for the read/write pattern, data size, and access frequency.

Relationship Modeling



Types of Relationships

Type	Example	Strategy
One-to-One	User ↔ Profile	Embed or Reference
One-to-Many	BlogPost ↔ Comments	Embed or Reference
Many-to-Many	Users ↔ Roles / Products ↔ Tags	Reference only

Relationship Modeling



- How to implement relationships in mongoose? :
 - There are two ways of implementing relationships:
 - Embedding:
 - storing related data inside the same document, as a sub-document or an array of sub-documents.
 - Referencing:
 - storing related data in separate documents, and connecting them using ObjectIds references.

Embedding



- How does embedding documents work :
 - Instead of creating separate collections (tables in a relational database) for related entities and linking them with references (like foreign keys), you embed the "child" document(s) directly within the "parent" document.

```
1 {
2   "_id": "u1",
3   "name": "Ali",
4   "email": "ali@example.com",
5   "profile": {
6     "bio": "Backend developer",
7     "avatar": "avatar.jpg"
8   }
9 }
```

Embedding



- **Benefits of Embedding:**
 - **Faster reads:** One document = one read from DB
 - **Simpler data model:** Easy to understand and access
 - **Atomic updates:** Updating the full document is safe
 - **Less joins/code:** No `.populate()` or `$lookup` needed

Embedding



- **When to Embed documents:**
 - When the embedded data is only used with the parent
 - When the data is not reused elsewhere
 - When you don't need to query the embedded data directly
 - When the embedded array is not very large (keep under ~100 subdocs ideally)
- **Avoid Embedding When:**
 - The sub-document grows indefinitely (e.g. chat messages)
 - You need to query/update sub-documents independently
 - The sub-document is shared across multiple parents

Referencing



- Referencing means storing related data in separate documents, and connecting them using ObjectID references.
- Think of it like creating a "foreign key" in relational databases.



```
1  {  
2    "_id": "u1",  
3    "name": "Ali",  
4    "email": "ali@example.com",  
5    "profile": "p1" // Reference to Profile._id  
6  }  
7
```

Referencing



- **When to Use Referencing**
 - **Use referencing when:**
 - The related data grows large (e.g., comments, orders)
 - You need to query/update sub-data independently
 - The related data is shared across multiple documents
 - You want to normalize your data
 - **Avoid referencing when:**
 - The data is always used together
 - You want fast, single-document reads

One-to-One Relationships



- **What Is a One-to-One Relationship?**
 - A One-to-One (1:1) relationship means that:
 - One document in a collection is linked to exactly one document in another collection.
- For example:
 - **A User has one Profile**
 - A Car has one Engine
 - **A Company has one Address**
 - **So, for each User → there is only one Profile.**

One-to-One Relationships



- **When Should You Use a One-to-One Relationship?**
 - Use a one-to-one relationship when:
 - You want to separate optional or sensitive information
 - The second document might grow large
 - You want to modularize your data (store it in its own collection)

Embedding

How to implement 1:1 embedding with mongoose :



```
1  const ProfileSchema = new mongoose.Schema({
2    bio: String,
3    avatar: String
4  });
5
6  const UserSchema = new mongoose.Schema({
7    name: String,
8    email: String,
9    profile: ProfileSchema
10  });
11
12  const User = mongoose.model('User', UserSchema);
13  export default User;
```



Embedding

Then you can create users with profiles using:



```
1  await User.create({
2    name: 'Ali',
3    email: 'ali@example.com',
4    profile: {
5      bio: 'Node.js backend developer',
6      avatar: 'ali.jpg'
7    }
8  });
9
10 const user = await User.findOne({ email: 'ali@example.com' });
11 console.log(user.profile.bio);
12
```

Referencing

How to implement 1:1 referencing with mongoose :



```
1  const ProfileSchema = new mongoose.Schema({
2    bio: String,
3    avatar: String,
4    social: {
5      twitter: String,
6      github: String
7    }
8  });
9  const UserSchema = new mongoose.Schema({
10    name: String,
11    email: String,
12    profile: {
13      type: mongoose.Schema.Types.ObjectId, ref: 'Profile'
14    }
15  });
```



Referencing

Then you can create a user profile and link it to the user using:



```
1  const profile = await Profile.create({
2    bio: 'I love Node.js!',
3    avatar: 'avatar.png',
4    social: {
5      twitter: '@mamdouh',
6      github: 'mamdouhdev'
7    }
8  });
9
10 const user = await User.create({
11   name: 'Mamdouh',
12   email: 'mamdouh@example.com',
13   profile: profile._id
14 });
```



Referencing

Now when we want to retrieve a user with their profile, we use the `.populate()` method to join tables:



```
1 console.log(await User.findById(user._id).populate('profile'));
2 {
3   "_id": "...",
4   "name": "Mamdouh",
5   "email": "mamdouh@example.com",
6   "profile": {
7     "_id": "...",
8     "bio": "I love Node.js!",
9     "avatar": "avatar.png",
10    "social": {
11      "twitter": "@mamdouh",
12      "github": "mamdouhdev"
13    }
14  }
15 }
```



mongoDB

One-to-Many Relationships



- **What Is a One-to-Many Relationship?**
 - A One-to-Many relationship means that one document is related to many other documents.
 - Real-world examples:
 - One user has many orders
 - One post has many comments
 - One category has many products
 - So, 1 "parent" → many "children"

One-to-Many Relationships



- **When Should You Use a One-to-Many Relationship?**
 - Use it when:
 - You have a clear “parent-child” relationship
 - The "many" side is logically grouped with the "one"
 - You want to store related but independent documents

Embedding

How to implement 1:M embedding with mongoose :



```
1  const CommentSchema = new mongoose.Schema({
2    text: String,
3    user: String
4  });
5
6  const PostSchema = new mongoose.Schema({
7    title: String,
8    content: String,
9    comments: [CommentSchema]
10  });
11
12  const Post = mongoose.model('Post', PostSchema);
```



mongoDB

Embedding

Then you can create comments for a post using:



```
1 // Find the post
2 const post = await Post.findById(post._id);
3
4 // Add new comment
5 post.comments.push({
6   text: 'Very helpful!',
7   user: 'Mamdouh'
8 });
9
10 // Save the updated post
11 await post.save();
```



mongoDB

Referencing

How to implement 1:M referencing with mongoose :



```
1  const PostSchema = new mongoose.Schema({
2    title: String,
3    content: String
4  });
5
6  const Post = mongoose.model('Post', PostSchema);
7
8  const CommentSchema = new mongoose.Schema({
9    text: String,
10   user: String,
11   postId: { type: mongoose.Schema.Types.ObjectId, ref: 'Post' }
12 });
13
14 const Comment = mongoose.model('Comment', CommentSchema);
```



mongoDB

Referencing

Then you can create comments for a post using:



```
1 // 1. Create a post
2 const post = await Post.create({
3   title: 'Mongo Relationships',
4   content: 'Let's learn about 1:M referencing!'
5 });
6
7 // 2. Create comments for that post (referencing it)
8 const comment1 = await Comment.create({
9   text: 'This is very helpful!',
10  user: 'Sara',
11  postId: post._id
12 });
```



mongoDB

Referencing

Finally, you could retrieve all post comments using:



```
1  const { text } = require("express");
2
3  const comments = await Comment.find({ postId: post._id });
4  console.log(comments);
5
6  // Output:
7
8  [
9    {
10      _id: '60c72b2f9b1e8c001c8e4d5a',
11      postId: '60c72b2f9b1e8c001c8e4d5a',
12      user: 'Sara',
13      text: 'this is very helpful'
14    },
15    ....
16  ]
```



mongoDB

Many-to-Many Relationships



- **What Is a Many-to-Many Relationship?**
 - A Many-to-Many relationship means:
 - A book can belong to many categories.
 - A category can contain many books.
 - This is different from One-to-Many (1:M) where only one side holds multiple references.

Many-to-Many Relationships



- **When Should You Use Many-to-Many?**
 - You should use M:M relationships when:
 - You need cross-linking between entities.
 - Both sides need to be aware of each other.
 - The relationship is flexible and dynamic (a book may be in several genres, and genres may change over time).
- **Real-world Examples:**
 - Books and Categories
 - Students and Courses
 - Users and Roles
 - Posts and Tags

Referencing

- You can't embed schemas recursively because it leads to a circular dependency
 - We can only use referencing if we want true, bi-directional M:M relationship:



```
1  const CategorySchema = new mongoose.Schema({
2    name: String,
3    books: [{
4      type: mongoose.Schema.Types.ObjectId,
5      ref: 'Book'
6    }]
7  });
8
9  const BookSchema = new mongoose.Schema({
10   title: String,
11   author: String,
12   categories: [{
13     type: mongoose.Schema.Types.ObjectId,
14     ref: 'Category'
15   }]
16  });
```



Referencing

Then you can create books and categories bi-directionally using:



```
1  const fiction = await Category.create({ name: 'Fiction' });
2  const history = await Category.create({ name: 'History' });
3
4  const book = await Book.create({
5    title: 'World Stories',
6    author: 'John Smith',
7    categories: [fiction._id, history._id]
8  });
```



Referencing

But because MongoDB is a non-relational database, we have to update our categories to include the new books too :



```
1  const fiction = await Category.create({ name: 'Fiction' });
2  const history = await Category.create({ name: 'History' });
3
4  const book = await Book.create({
5    title: 'World Stories',
6    author: 'John Smith',
7    categories: [fiction._id, history._id]
8  });
9
10 fiction.books.push(book._id);
11 history.books.push(book._id);
12 await fiction.save();
13 await history.save();
```



Populating Responses



- What Is `.populate()`?
 - `.populate()` is a Mongoose method used to automatically replace referenced ObjectIds in a document with the actual documents from the referenced collection.
 - This is a crucial feature for handling relationships between data in MongoDB, where relationships are often represented by storing the `_id` of a related document.

populate

- How it works:
 1. Referencing: In your Mongoose schema, you define a field that references another model using `type: Schema.Types.ObjectId` and the `ref` option, which specifies the name of the referenced model.



```
1  const storySchema = new mongoose.Schema({  
2    author: { type: mongoose.Schema.Types.ObjectId, ref: 'Person' },  
3    title: String,  
4  });
```


populate

- How it works:
 1. Query: When you query for documents from the first collection, you can use `.populate()` on the query to replace the `ObjectId` in the referenced field with the actual document data from the related collection.
 2. Population: When you query for documents from the first collection, you can use `.populate()` on the query to replace the `ObjectId` in the referenced field with the actual document data from the related collection.



```
1 Story.findOne({ title: 'My Awesome Story' })
2   .populate('author')
3   .exec((err, story) => {
4     if (err) return handleError(err);
5     console.log('The author is %s', story.author.name);
6   });
```

Populating Responses



- **Key features and uses:**
 - **Simplifies data retrieval:**
 - Eliminates the need for manual lookups or multiple queries to fetch related data.
 - **Populate single or multiple documents:**
 - Can be used to populate a single field or an array of referenced documents.
 - **Populate multiple paths:**
 - You can chain `.populate()` calls to populate multiple fields in a single query.

Populating Responses



- **Key features and uses:**
 - **Select specific fields:**
 - You can specify which fields from the populated document you want to include using the select option.
 - **Query conditions within population:**
 - You can apply conditions and options to the populated documents themselves, such as limiting the number of sub-documents populated from an array.

Populating Responses



Parameters

Parameter	Type	Description
path	String	Field to populate (must be a reference)
select	String or Object	Fields to include/exclude ('name email' or '-password')
model	String	Explicit model name (optional)
match	Object	Filter criteria for populated documents
options	Object	Pagination/sorting: { limit, skip, sort }

Populating Responses



- Examples:
 - Basic populate: `Post.find().populate('author');`
 - Select specific fields: `Post.find().populate('author', 'name email -_id');`
 - Filtering referenced documents:

```
Post.find().populate({ path: 'author', match: { name: /john/i } });
```

- Sorting and limiting populated documents:

```
Post.find().populate({ path: 'comments', options: { sort: { createdAt: -1 }, limit: 5 } });
```

Aggregation Pipeline



Aggregation Pipeline



- **What Is the Aggregation Pipeline?**
 - The Aggregation Pipeline in MongoDB is a framework to process data records, transform them, and return computed results — all within the database, not in your app.
 - Each stage performs a specific operation, like filtering, grouping, or sorting, and the output of one stage becomes the input for the next.
 - This framework provides a powerful and flexible way to work with data in MongoDB.

Aggregation Pipeline



- **Core Concepts:**
 - **Stages:**
 - The pipeline is composed of multiple stages, each performing a specific operation on the documents.
 - **Data Flow:**
 - Documents enter the pipeline and are transformed by each stage in sequence.
 - **Result:**
 - The final output of the pipeline is the result of all the operations performed on the input documents.

Aggregation Stages



- `$match`: Filters documents based on specified criteria.
- `$group`: Groups documents based on a key and performs calculations on the groups.
- `$sort`: Sorts documents based on specified fields.
- `$project`: Reshapes the structure of documents by including or excluding fields, or adding new fields.
- `$unwind`: Deconstructs an array field from an input document into multiple documents, one for each element in the array.

Aggregation Stages



- `$limit`: Limits the number of documents passed to the next stage.
- `$skip`: Skips a specified number of documents before passing the remaining ones to the next stage.
- `$lookup`: Performs a left outer join with another collection.
- `$addFields`: Adds new fields to documents.
- `$out`: Writes the results of the aggregation to a new collection.

\$sort



- What It Does:

- \$sort sorts the documents that pass through the aggregation pipeline based on specified fields, either in ascending (1) or descending (-1) order.
- Equivalent to **ORDER BY** price **DESC** in **SQL**



```
1 { name: "Book A", price: 20 }
2 { name: "Book B", price: 10 }
3 { name: "Book C", price: 30 }
4
5 db.products.aggregate([
6   { $sort: { price: -1 } }
7 ]);
```

Output



```
1 // output
2 { name: "Book C", price: 30 }
3 { name: "Book A", price: 20 }
4 { name: "Book B", price: 10 }
```

Aggregation Stages



- Tips & Notes:
 - **\$sort** must appear after **\$match**, **\$project**, or **\$group** if you're using them. sort operates on the current pipeline result.
 - Sorting large datasets may impact performance if no proper index exists.
 - Can sort by multiple fields: **{ price: 1, name: -1 }**.

\$group



mongoDB

- \$group groups documents by a specified key and lets you perform aggregate operations like:
 - sum, avg, min, max
 - push, addToSet
 - first, last, etc.
- Each output document represents one group.

```
1 // Sample data for the sales collection
2 { item: "Book", quantity: 5, price: 10 }
3 { item: "Pen", quantity: 10, price: 2 }
4 { item: "Book", quantity: 2, price: 10 }
5
6 db.sales.aggregate([
7   {
8     $group: {
9       _id: "$item",
10      totalQuantity: { $sum: "$quantity" }
11    }
12  }
13 ]);
```

Output

```
1 // Output:
2 { _id: "Book", totalQuantity: 7 }
3 { _id: "Pen", totalQuantity: 10 }
```

Aggregation Stages



Grouping Common Accumulators

Accumulator	What It Does
\$sum	Total sum of values
\$avg	Average value
\$min, \$max	Minimum or maximum value
\$first, \$last	First/last value by sort order
\$push	Adds values to array (duplicates ok)
\$addToSet	Adds values to array (unique only)

Aggregation Stages



- **Tips & Notes:**

- The `_id` is required, it defines how you group.
 - Use `null` if you want one group for everything.
- You can group by multiple fields using an object:

`_id: { item: "$item", category: "$category" }`

- Use `$project` before `$group` if you want to reshape documents first.
- SQL Equivalent: `SELECT item, SUM(quantity)`
`FROM sales`
`GROUP BY item;`

\$project



- \$project allows you to reshape a document by:
 - Including or excluding specific fields.
 - Renaming fields.
 - Adding new fields by computing values.
 - Control document structure for the next pipeline stage.
 - SQL Equivalent of: **SELECT** salary * 1.1 **AS** bonusSalary **FROM** employees

```
1 {
2   "_id": ObjectId("..."),
3   "firstName": "John",
4   "lastName": "Doe",
5   "age": 60,
6   "salary": 5000
7 }
8 db.employees.aggregate([
9   {
10    $project: {
11      _id: 0,
12      name: 1,
13      age: 1,
14      isSenior: { $gte: ["$age", 50] },
15      fullName: { $concat: ["$firstName", " ", "$lastName"] },
16      bonusSalary: { $multiply: ["$salary", 1.1] }
17    }
```

Output

```
1 {
2   "name": "John",
3   "age": 60,
4   "isSenior": true,
5   "fullName": "John Doe",
6   "bonusSalary": 5500
7 }
```


Aggregation Stages



Common **\$project** Operators

Operator	Description
1 / 0	Include (1) or exclude (0) a field
\$add, \$subtract, \$multiply, \$divide	Math operations
\$concat	Join strings "concatenate"
\$toUpper, \$toLower	Change string case
\$gte, \$lte, \$eq, \$cond	Logical comparisons and conditions
\$arrayElemAt	Get specific index from array
\$dateToString	Format date fields

- Takes an array field from input documents and outputs a separate document for each element in the array.

```
1  const employees = [  
2    { name: "Alice", skills: ["JavaScript", "React"]  
3    { name: "David", department: "Design" },  
4  ];  
5  
6  Employee.aggregate([  
7    {  
8      $unwind: {  
9        path: "$skills",  
10       includeArrayIndex: "skillIndex",  
11       preserveNullAndEmptyArrays: false  
12     }  
13   }  
14 ]));
```

Output →

```
// output  
[  
  {  
    _id: new ObjectId('687b0dae9aba05bc108b4e8e'),  
    name: 'Alice',  
    skills: 'JavaScript',  
    department: 'Engineering',  
    __v: 0,  
    skillIndex: 0  
  },  
  {  
    _id: new ObjectId('687b0dae9aba05bc108b4e8e'),  
    name: 'Alice',  
    skills: 'React',  
    department: 'Engineering',  
    __v: 0,  
    skillIndex: 1  
  },  
]
```

\$limit

- The **\$limit** stage limits the number of documents that pass through the aggregation pipeline to a specified number.
- Usually comes after sorting. Ex. Limit for the top five most expensive items

```
1  await Product.insertMany([
2    { name: 'Laptop', price: 1200 },
3    { name: 'Smartphone', price: 800 },
4    { name: 'Monitor', price: 300 },
5    { name: 'Keyboard', price: 100 },
6    { name: 'Tablet', price: 400 },
7    { name: 'Smartwatch', price: 250 },
8    { name: 'Desk Lamp', price: 60 },
9  ]);
10
11  const topProducts = await Product.aggregate([
12    { $sort: { price: -1 } },
13    { $limit: 5 }
14  ]);
```

Output

Top 5 Most Expensive Products:

```
[
  {
    _id: new ObjectId('687c45df56d122b046911d6d'),
    name: 'Laptop',
    price: 1200,
    __v: 0
  },
  {
    _id: new ObjectId('687c45df56d122b046911d6e'),
    name: 'Smartphone',
    price: 800,
    __v: 0
  },
  {
    _id: new ObjectId('687c45df56d122b046911d71'),
    name: 'Tablet',
    price: 400,
    __v: 0
  },
  {
    _id: new ObjectId('687c45df56d122b046911d6f'),
    name: 'Monitor',
    price: 300,
    __v: 0
  },
  {
    _id: new ObjectId('687c45df56d122b046911d72'),
    name: 'Smartwatch',
    price: 250,
    __v: 0
  }
]
```



\$skip

- \$skip is a stage in the MongoDB Aggregation Pipeline that skips a specified number of documents from the input and passes the remaining documents to the next stage.
- Think of it like the SQL **OFFSET** keyword or **LIMIT** x **OFFSET** y

```
1  await Product.insertMany([
2    { name: 'Laptop', price: 1200 },
3    { name: 'Phone', price: 800 },
4    { name: 'Tablet', price: 600 },
5    { name: 'Monitor', price: 300 },
6    { name: 'Keyboard', price: 100 },
7    { name: 'Mouse', price: 50 },
8    { name: 'Printer', price: 400 },
9    { name: 'Speaker', price: 200 },
10   { name: 'Webcam', price: 150 },
11   { name: 'Microphone', price: 250 }
12  ]);
13
14  const result = await Product.aggregate([
15    { $sort: { price: -1 } },
16    { $skip: 5 },
17    { $limit: 5 }
18  ]);
```

Output

```
Products after skipping top 5 expensive ones:
[
  {
    _id: new ObjectId('687c49221e4da9d4c480b2b5'),
    name: 'Microphone',
    price: 250,
    __v: 0
  },
  {
    _id: new ObjectId('687c49221e4da9d4c480b2b3'),
    name: 'Speaker',
    price: 200,
    __v: 0
  },
  {
    _id: new ObjectId('687c49221e4da9d4c480b2b4'),
    name: 'Webcam',
    price: 150,
    __v: 0
  },
  {
    _id: new ObjectId('687c49221e4da9d4c480b2b0'),
    name: 'Keyboard',
    price: 100,
    __v: 0
  },
  {
    _id: new ObjectId('687c49221e4da9d4c480b2b1'),
    name: 'Mouse',
    price: 50,
    __v: 0
  }
]
```

\$lookup



- \$lookup allows you to join data from another collection, similar to JOIN in SQL.
- It performs a left outer join, which means:
 - All documents from the “left” collection (main one) will be returned.
 - Matching documents from the “right” collection (foreign) will be included if they exist.
 - If no match is found, an empty array is returned in the field.

\$lookup



Important Notes:

- from must be the collection name, not the model name (e.g., 'courses', not 'Course')
- Only works across collections in the same database
- **\$lookup** adds an array, even if only one document is matched
- Combine with **\$unwind** if you want to flatten single-value arrays
- Cannot populate using **\$lookup**; use one or the other
- Heavy **\$lookups** on large datasets can be slow – use wisely

\$lookup



\$lookup vs .populate() vs SQL JOIN

Feature	\$lookup	.populate()	SQL JOIN
Where it runs	DB-level (fast)	App-level (JS)	DB-level (fast)
Returns	Array (always)	Object or Array	Rows
Flexibility	Very high (can filter, reshape, etc.)	Low (automatic only)	High
Use case	Reports, stats, dashboards	Normal app behavior	Complex queries

```

await Student.insertMany([
  { name: 'Alice', courseIds: [course1._id, course2._id] },
  { name: 'Charlie', courseIds: [] }
]);

const result = await Student.aggregate([
  {
    $lookup: {
      from: 'courses',
      localField: 'courseIds',
      foreignField: '_id',
      as: 'enrolledCourses'
    },
  },
  {
    $project: {
      name: 1,
      enrolledCourses: 1
    }
  }
]);

```

Output

```

1  [
2    {
3      "_id": "687c4b81c5492eef52b698f2",
4      "name": "Alice",
5      "courseIds": [
6        "687c4b81c5492eef52b698ec",
7        "687c4b81c5492eef52b698ee"
8      ],
9      "__v": 0,
10     "enrolledCourses": [
11       {
12         "_id": "687c4b81c5492eef52b698ec",
13         "title": "Math",
14         "__v": 0
15       },
16       {
17         "_id": "687c4b81c5492eef52b698ee",
18         "title": "Physics",
19         "__v": 0
20       }
21     ]
22   },
23   {
24     "_id": "687c4b81c5492eef52b698f4",
25     "name": "Charlie",
26     "courseIds": [],
27     "__v": 0,
28     "enrolledCourses": []
29   }
30 ]
31

```


\$addFields



- **\$addFields** adds new fields to each document in the pipeline.
- It can also modify existing fields.
- It's non-destructive; Existing fields stay unless you overwrite them.
- It's like adding a computed column in a SELECT query:

```
SELECT name, salary, salary * 0.1 AS bonus FROM employees;
```

\$addFields



Use Cases:

- Computed fields (discounts, totals, etc.)
- Text transformations (**\$toUpper**, **\$concat**)
- Extracting from arrays (**\$arrayElemAt**)
- Combining with **\$project** for cleaner output
- Don't use it if you only want to remove or rename fields; Use **\$project** for that.



mongoDB

\$addFields vs \$project

Feature	\$addFields	\$project
Adds new fields	Yes	Yes
Removes fields	No (keeps all existing fields)	Yes (explicitly include/exclude fields)
Modifies fields	Yes	Yes (can override fields with new expressions)
Default behavior	Keeps all existing fields	Excludes all unless you include them manually
Best used for	Calculating and appending new data	Shaping final output (like SELECT in SQL)

```

await Product.insertMany([
  { name: 'Laptop', price: 1000, tags: ['electronics', 'computing'] },
  { name: 'Phone', price: 500, tags: ['electronics', 'mobile'] },
  { name: 'Book', price: 40, tags: ['reading', 'education'] }
]);

const result = await Product.aggregate([
  {
    $addFields: {
      discountPrice: { $multiply: ['$price', 0.8] },
      upperCaseName: { $toUpper: '$name' },
      firstTag: { $arrayElemAt: ['$tags', 0] }
    }
  },
  {
    $project: {
      _id: 0,
      name: 1,
      price: 1,
      discountPrice: 1,
      upperCaseName: 1,
      firstTag: 1
    }
  }
]);

```

Output

```

1 //With Added Fields:
2 [
3   {
4     name: 'Laptop',
5     price: 1000,
6     discountPrice: 800,
7     upperCaseName: 'LAPTOP',
8     firstTag: 'electronics'
9   },
10  {
11    name: 'Phone',
12    price: 500,
13    discountPrice: 400,
14    upperCaseName: 'PHONE',
15    firstTag: 'electronics'
16  },
17  {
18    name: 'Book',
19    price: 40,
20    discountPrice: 32,
21    upperCaseName: 'BOOK',
22    firstTag: 'reading'
23  }
24 ]
25

```

\$out



- **\$out** writes the final result of the aggregation pipeline to a collection (either replacing it or creating a new one).
- Think of it like saving the result of a complex query into its own collection.
- Must be the last stage in the pipeline.
- It will replace the contents of the target collection.
- When to use **\$out**:
 - Caching expensive aggregation results.
 - Generating reports.
 - Archiving filtered data.
 - Creating flattened/denormalized views of data.

```

1  await Product.insertMany([
2    { name: 'Laptop', price: 1000, tags: ['electronics', 'computing']
3    },
4    { name: 'Phone', price: 500, tags: ['electronics', 'mobile'] },
5    { name: 'Book', price: 40, tags: ['reading', 'education'] }
6  ]);
7  await Product.aggregate([
8    {
9      $addFields: {
10        discountPrice: { $multiply: ['$price', 0.8] },
11        upperCaseName: { $toUpper: '$name' },
12        firstTag: { $arrayElemAt: ['$tags', 0] }
13      }
14    },
15    {
16      $project: {
17        _id: 0,
18        name: 1,
19        price: 1,
20        discountPrice: 1,
21        upperCaseName: 1,
22        firstTag: 1
23      }
24    },
25    {
26      $out: 'discounted_products'
27    }
28  ]);

```

Output

```

// Discounted Products:
[
  {
    _id: new ObjectId('687c58182ca71a8e723d25c3'),
    name: 'Laptop',
    price: 1000,
    discountPrice: 800,
    upperCaseName: 'LAPTOP',
    firstTag: 'electronics'
  },
  {
    _id: new ObjectId('687c58182ca71a8e723d25c4'),
    name: 'Phone',
    price: 500,
    discountPrice: 400,
    upperCaseName: 'PHONE',
    firstTag: 'electronics'
  },
  {
    _id: new ObjectId('687c58182ca71a8e723d25c5'),
    name: 'Book',
    price: 40,
    discountPrice: 32,
    upperCaseName: 'BOOK',
    firstTag: 'reading'
  }
]

```

\$merge



- **\$merge** writes the output of an aggregation pipeline to a collection, just like \$out, but with more control over what happens when documents with matching _id already exist.
- Can output to a collection in the same or different database.
- Can output to the same collection that is being aggregated.
- Creates a new collection if the output collection does not already exist.
- Can incorporate results (insert new documents, merge documents, replace documents, keep existing documents)

```

1  await Product.aggregate([
2    {
3      $match: { hasDiscount: true }
4    },
5    {
6      $addFields: {
7        discountPrice: {
8          $subtract: [
9            '$price',
10           { $multiply: ['$price', '$discountRate'] }
11         ]
12       },
13       uppercaseName: { $toUpper: '$name' },
14       firstTag: { $arrayElemAt: ['$tags', 0] }
15     },
16   },
17   {
18     $project: {
19       _id: 1,
20       name: 1,
21       price: 1,
22       discountRate: {
23         $concat: [
24           { $toString: { $multiply: [100, '$discountRate'] } }, '%'
25         ],
26         discountPrice: 1,
27         uppercaseName: 1,
28         firstTag: 1
29       }
30     },
31   },
32   {
33     $merge: {
34       into: 'discounted_products',
35       whenMatched: 'merge',
36       whenNotMatched: 'insert'
37     }
38   }
39 ]);

```

Output

```

//Discounted Products:
[
  {
    _id: new ObjectId('687c6553e5f8e8623dc7a5fb'),
    discountPrice: 800,
    discountRate: '20%',
    firstTag: 'electronics',
    name: 'Laptop',
    price: 1000,
    uppercaseName: 'LAPTOP'
  },
  {
    _id: new ObjectId('687c6553e5f8e8623dc7a5fd'),
    discountPrice: 36,
    discountRate: '10%',
    firstTag: 'education',
    name: 'Book',
    price: 40,
    uppercaseName: 'BOOK'
  },
  {
    _id: new ObjectId('687c6553e5f8e8623dc7a5fe'),
    discountPrice: 170,
    discountRate: '15%',
    firstTag: 'furniture',
    name: 'Desk',
    price: 200,
    uppercaseName: 'DESK'
  }
]

```